

# Deep Learning

## Lab 5: Regularization

DataLab, 2022

Department of Computer Science,  
National Tsing Hua University, Taiwan

# Regularization

techniques that improve the  
**generalizability** of a trained model

# Outline

- Scikit-learn
- Learning Theory
  - Error Curves and Model Complexity
  - Learning Curves and Sample Complexity
- Weight Decay
  - Ridge Regression
  - LASSO
- Validation
- Assignment

# Outline

- Scikit-learn
- Learning Theory
  - Error Curves and Model Complexity
  - Learning Curves and Sample Complexity
- Weight Decay
  - Ridge Regression
  - LASSO
- Validation
- Assignment

# Scikit-learn

- Scikit-learn is a free software machine learning library for the Python programming language
- It features various classification, regression and clustering algorithms
  - including SVM (support vector machines), Random Forests, gradient boosting, k-means and DBSCAN, and is designed to interoperate with the Python numerical and scientific libraries NumPy and SciPy
- `pip install scikit-learn` / `conda install scikit-learn`

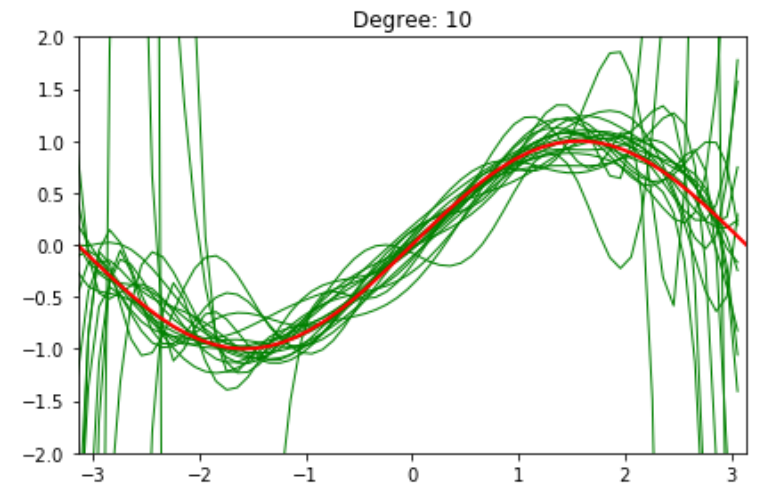




# Outline

- Scikit-learn
- Learning Theory
  - Error Curves and Model Complexity
  - Learning Curves and Sample Complexity
- Weight Decay
  - Ridge Regression
  - LASSO
- Validation
- Assignment



# Learning Theory

- Learning theory provides a means to understand the generalizability of the model
- **Model complexity** plays a crucial role
  - Too simple: high bias and underfitting
  - Too complex: high variance and overfitting





# Error Curves and Model Complexity

- It is relatively hard to observe the figures showed in the last slide, since normally we will never know the data distribution of ground truth (red line in the last slide)
- Instead, we can get those information by observing the training and testing error curve







# Double Descent Curves in **Modern** Machine Learning\*\*



Reconciling modern machine learning practice and the bias-variance trade-of (PNAS'19)  
Double-descent curves in neural networks: a new perspective using Gaussian processes (arXiv'21)


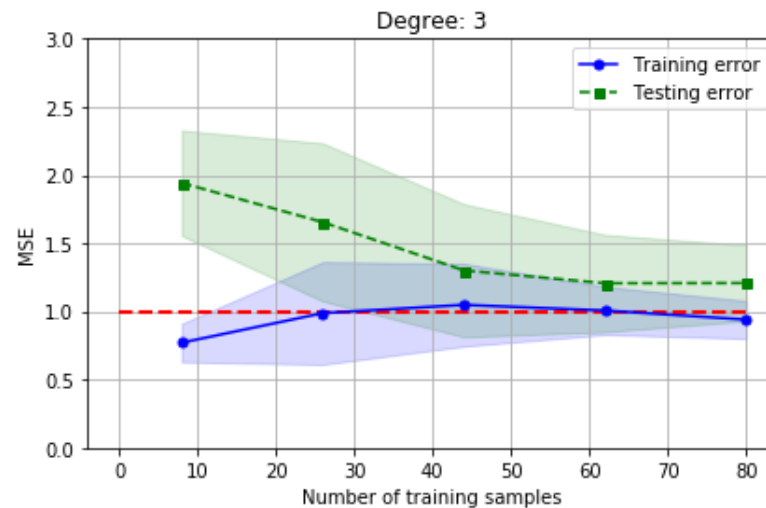

# Error Curves and Model Complexity

- Although the error curve visualizes the impact of model complexity, the bias-variance tradeoff holds **only when you have sufficient training examples**



# Learning Curves and Sample Complexity

- The bounding methods of learning theory tell us that a model is likely to overfit regardless of its complexity **when the size of training set is small**. The **learning curves** are a useful tool for understanding how much training examples are sufficient



# Outline

- Scikit-learn
- Learning Theory
  - Error Curves and Model Complexity
  - Learning Curves and Sample Complexity
- **Weight Decay**
  - Ridge Regression
  - LASSO
- Validation
- Assignment

# Weight Decay

- A common regularization approach. The idea is to add a term in the cost function against complexity

- Ridge Regression ( $L_2$ )

$$\arg \min_{\mathbf{w}, b} \|\mathbf{y} - (X\mathbf{w} - b\mathbf{1})\|^2 + \alpha \|\mathbf{w}\|^2$$

- LASSO ( $L_1$ )

$$\arg \min_{\mathbf{w}, b} \|\mathbf{y} - (X\mathbf{w} - b\mathbf{1})\|^2 + \alpha \|\mathbf{w}\|_1$$

# Ridge Regression

- A small value  $\alpha$  drastically reduces the testing error. Nevertheless, it's not a good idea to increase  $\alpha$  forever, since it will over-shrink the coefficients of  $w$  and result in underfitting

$$\arg \min_{w,b} \|y - (Xw - b\mathbf{1})\|^2 + \alpha \|w\|^2$$

```
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error

poly = PolynomialFeatures(degree=3)
X_poly = poly.fit_transform(X_std)
X_train, X_test, y_train, y_test = train_test_split(
    X_poly, y, test_size=0.3, random_state=0)

for a in [0, 1, 10, 100, 1000]:
    lr_rg = Ridge(alpha=a)
    lr_rg.fit(X_train, y_train)

    y_train_pred = lr_rg.predict(X_train)
    y_test_pred = lr_rg.predict(X_test)

    print('\n[Alpha = %d]' % a )
    print('MSE train: %.2f, test: %.2f' % (
        mean_squared_error(y_train, y_train_pred),
        mean_squared_error(y_test, y_test_pred)))
```

```
[Alpha = 0]
MSE train: 0.00, test: 19958.68

[Alpha = 1]
MSE train: 0.73, test: 23.05

[Alpha = 10]
MSE train: 1.66, test: 16.83

[Alpha = 100]
MSE train: 3.60, test: 15.16

[Alpha = 1000]
MSE train: 8.81, test: 19.22
```

# LASSO

- An alternative weight decay approach that can lead to sparse  $w$  is the LASSO. Depending on the value of  $\alpha$ , certain weights can become zero much faster than others

$$\arg \min_{w,b} \|y - (Xw - b\mathbf{1})\|^2 + \alpha \|w\|_1$$

```
from sklearn.linear_model import Lasso
from sklearn.metrics import mean_squared_error

for a in [0.001, 0.01, 0.1, 1, 10]:
    lr_rg = Lasso(alpha=a)
    lr_rg.fit(X_train, y_train)

    y_train_pred = lr_rg.predict(X_train)
    y_test_pred = lr_rg.predict(X_test)

    print('\n[Alpha = %.2f]' % a )
    print('MSE train: %.2f, test: %.2f' % (
        mean_squared_error(y_train, y_train_pred),
        mean_squared_error(y_test, y_test_pred)))
```

[Alpha = 0.0000]  
MSE train: 0.55, test: 61.02

[Alpha = 0.0010]  
MSE train: 0.64, test: 29.11

[Alpha = 0.0100]  
MSE train: 1.52, test: 19.51

[Alpha = 0.1000]  
MSE train: 4.34, test: 15.52

[Alpha = 1.0000]  
MSE train: 14.33, test: 22.42

[Alpha = 10.0000]  
MSE train: 55.79, test: 53.42



# Ridge vs LASSO

- Why is LASSO sparse?

Ridge: [0.5, 0.5, 0.5, 0.5]

Initial weights: [1, 0.5, 1, 0.5]

LASSO: [0.5, 0, 0.5, 0]






# Ridge vs LASSO

- Why is LASSO sparse?


$$\arg \min_{\mathbf{w}, b} \frac{1}{2N} \|\mathbf{y} - (\mathbf{X}\mathbf{w} - b\mathbf{1})\|^2 + \alpha \|\mathbf{w}\|_1$$

- The surface of the cost function is the sum of SSE (blue contours) and 1-norm (red contours)
- Optimal point locates on some axes



# Ridge vs LASSO

- LASSO can also be treated as a supervised **feature selection** technique when choosing a suitable regularization strength  $\alpha$  to make only part of coefficients become exactly zeros




# Outline

- Scikit-learn
- Learning Theory
  - Error Curves and Model Complexity
  - Learning Curves and Sample Complexity
- Weight Decay
  - Ridge Regression
  - LASSO
- Validation
- Assignment

# Validation

- Another useful regularization technique that helps us decide the proper value of **hyperparameters**
- The idea is to split your data into the training, validation, and testing sets and then select the best value based on validation performance
- NOTE: It is important that we should never peek testing data during training



# Validation

[Degree = 1]  
MSE train: 25.00, valid: 21.43, test: 32.09

[Degree = 2]  
MSE train: 9.68, valid: 14.24, test: 20.24

[Degree = 3]  
MSE train: 3.38, valid: 17.74, test: 18.63

[Degree = 4]  
MSE train: 1.72, valid: 16.67, test: 30.98

[Degree = 5]  
MSE train: 0.97, valid: 59.73, test: 57.02


[Degree = 6]  
MSE train: 0.60, valid: 1444.08, test: 33189.41

# Outline

- Scikit-learn
- Learning Theory
  - Error Curves and Model Complexity
  - Learning Curves and Sample Complexity
- Weight Decay
  - Ridge Regression
  - LASSO
- Validation
- Assignment

# Assignment

- In this assignment, we would like to predict the success of shots made by basketball players in the NBA



# Assignment

- In this assignment, we would like to predict the success of shots made by basketball players in the NBA
  - **y\_test** is hidden this time
  - Allow to use **any model** you have learned before to achieve the best accuracy
  - Select the best **3 features**, and show the accuracy with only those
- Hint
  - **Preprocess the data** to help your training
  - Since you don't have y\_test this time, you may need to **split a validation set** for checking your performance
  - It is possible to use regression model as a classifier, for example [RidgeClassifier](#)



# Assignment

- Submit to **eeclash** with your:
  - **ipynb** (Lab05\_{student\_id}.ipynb)
  - **Prediction** (Lab05\_{student\_id}\_y\_pred.csv)
- The notebook should contain
  - How you **evaluate** your model
  - **All models** you have tried and the result
  - Plot the **error curve** of your best model and tell if it is **over-fit or not**
  - The **top-3 features** you find and how you find it
  - A **brief report** what you do in this assignment
- Deadline: **2022-10-06(Thur) 23:59**

# About Competition

- Students will group (2~4 people a group)
- This class suggests each group of students to **prepare a GPU card** to perform the necessary computing.
  - NO GPU CARD PROVIDED IN THE CLASS.
- Sign up for your group before the class on 10/4(Tue)

# About Competition

≡ kaggle

+ Create

🏠 Home

🏆 Competitions

📊 Datasets

<> Code

💬 Discussions

🎓 Learn

▼ More

📅 View Active Events

🔍 Search

Sign In

Register

🌟 Community Prediction Competition

## Datalab Cup4: Unlearnable Datasets CIFAR-10

Competition for CS565600 Deep Learning

17 teams · 8 months ago

Overview Data Code Discussion Leaderboard Rules

Join Competition

...

### Leaderboard






📄 Raw Data






















































🔄 Refresh

🔍 Search leaderboard

Public Private

This leaderboard is calculated with approximately 50% of the test data. The final results will be based on the other 50%, so the final standings may be different.

#	Team	Members	Score	Entries	Last	Code
1	我愛DL:)		0.87520	14	8mo	
2	電極史派克		0.87200	19	8mo	
3	開局就殘血		0.84240	17	8mo	
4	我們可能會退選 不要找我們組隊		0.83100	10	8mo	
5	Pytorture		0.82580	12	8mo	

1	▲ 1	電極史派克	   	0.87260	19	8mo
2	▲ 1	開局就殘血	   	0.84180	17	8mo
3	▲ 1	我們可能會退選 不要找我們組隊	 	0.84040	10	8mo
4	▲ 1	Pytorture	  	0.83080	12	8mo
5	▲ 2	百鬼組	   	0.82660	8	8mo
6	—	地獄兄弟	  	0.82260	13	8mo
 cifar_TA_80.csv				0.81780		
7	▼ 6	我愛DL:)	   	0.80660	14	8mo
8	—	kk	  	0.80360	7	8mo
9	▲ 1	Papaya	   	0.78640	8	8mo
10	▼ 1	Pytorch	   	0.78040	10	8mo
11	—	讓我活下來	   	0.74460	10	8mo
12	—	麻椅上訴	   	0.74080	10	8mo
13	—	歐  比	  	0.72400	12	8mo
14	▲ 1	哭啊	  	0.69880	11	8mo
15	▼ 1	Team123	   	0.68760	24	8mo
16	—	無敵的吧	  	0.68560	2	8mo
 cifar_TA_60.csv				0.63740		
17	—	深度度噲噲噲	   	0.41960	3	8mo